

How repeated reporting becomes one trustworthy story

Clustering reduces repetition without confusing two different events that merely sound alike. Each level adds intelligence while preserving the underlying evidence.

1

Deterministic normalization

Canonicalize URLs, remove tracking parameters, clean titles and punctuation, unify source identities, and collapse exact or nearly exact duplicates before expensive processing.

2

Semantic embeddings

Convert the title, RSS description and useful metadata into a numerical meaning vector. Differently worded reports can become candidates even when their headlines share few words.

3

Recent candidate comparison

Compare only plausible recent neighbors selected by time, category, geography, named entities and vector similarity. Never search the entire database for every article.

4

Consistency verification

Before merging, verify that people, organizations, location, event time, major numbers, topic and event nature agree. Similarity proposes the match; consistency decides it.

5

Source preservation

Keep every accepted article attached to the canonical cluster. The final story retains publisher attribution, direct links, evidence history and the ability to re-verify future updates.

The safety rule

A shared name or high vector score is never enough. Election rounds, court hearings, market moves and sports matches remain separate unless the event facts align. Thresholds are calibrated against a manually reviewed real-news test set.


How one shared news pool becomes a personal briefing

Gemini does not create a separate news universe for every subscriber. Stories receive structured metadata once; transparent code then filters, scores and diversifies them for each delivery.

One shared intelligence record per cluster



The result



No click tracking. No per-user AI summarization. Explicit choices remain in control.